vMCA: Memory Capacity Aggregation and Management in Cloud Environments

Luis A. Garrido
Barcelona Supercomputing Center
Barcelona, Spain
Email: luis.garrido@bsc.es

Paul Carpenter
Barcelona Supercomputing Center
Barcelona, Spain
Email: paul.carpenter@bsc.es

Abstract—In cloud infrastructures, the VMs within the computing nodes generate varying memory demand profiles. When memory utilization reaches its limits due to this demand, costly operations such as (virtual) disk accesses and/or VM migrations can occur. Since some nodes might have memory sitting idle, some of these costly operations could be avoided by making the idle memory available to the nodes that need it. Considering this, new architectures have been introduced that provide hardware support for a shared global address space that, together with fast interconnects, can share resources across nodes. Consequently, memory becomes a global resource.

This paper presents a memory capacity aggregation mechanism for cloud environments called vMCA (Virtualized Memory Capacity Aggregation) based on Xen's Transcendent Memory (Tmem). vMCA distributes the system's total memory within a single-node and globally across computing nodes using a user-space process with high-level memory management policies. We evaluate our solution using CloudSuite 3.0 on Linux and Xen. Our results demonstrate a peak running time improvement of 76.8% when aggregating memory, and of 37.5% when aggregating memory and implementing our memory management policies.

I. Introduction

Cloud data centres are currently built using "share nothing" servers, each provisioned with their own computing resources, communicating over a TCP/IP (or similar) network. New system architectures [7], [5], [8], however, have been proposed that present a shared global physical address space and use a fast interconnect to share physical resources through the memory hierarchy. An example is the UNIMEM memory model, which is an important feature of the "EuroEXA" family of projects [7], [12]. In such systems, remote memory is addressable at low latency using Remote DMA and/or load/store instructions. Cache coherency is only enforced inside a node (known as a *coherence island*), which avoids global coherence traffic and enables scalability to large numbers of nodes.

Such architectures present an opportunity to aggregate physical memory capacity. In cloud environments, memory is currently managed at the node level, with the hypervisor distributing this capacity among its Virtual Machines (VMs), each of which hosts a guest OS. The hypervisor usually overcommits its memory, which allows the node to have more active VMs by assuming they won't use their full allocation. This improves the utilization of the node's memory but also increases pressure on it, even more so if the VMs have varying demands for memory. To alleviate this pressure,

dynamic memory management mechanisms are implemented, such as memory ballooning and hotplug. Xen's Transcendent Memory (Tmem) [1] is another way to make additional memory capacity available to the VMs, which works by anticipatedly pooling all underutilized memory capacity.

This paper presents vMCA (virtualized Memory Capacity Aggregation), an extension of Xen's Tmem that leverages virtualization software and a global address space to allow the whole system's memory capacity to be shared among nodes. vMCA extends and improves upon an earlier proposal known as GV–Tmem [9], in order to handle resiliency, and with a deeper discussion and improvements of the memory aggregation and allocation policies. Like GV–Tmem, vMCA introduces minimal changes to the hypervisor to enforce constraints on local and remote page allocation. The complexity is situated in a user-space Memory Manager (MM) in the privileged domain of the nodes, which implements a resilient and reliable mechanism for distributing memory capacity across nodes according to a high-level memory management policy.

The main contributions of this paper are:

1) A software stack to aggregate memory capacity focused on resiliency and efficiency.
2) A multi-level user-space mechanism for allocation/management of aggregated memory capacity, with special considerations for memory allocation within a node and memory distribution across nodes.
3) A complete analysis of high-level policies for memory aggregation and allocation.

This paper is organized as follows. Section II provides the necessary background. Section III presents vMCA and its components, explaining its resiliency to failures and the necessary hardware support. Section IV describes the experimental methodology and Section V shows the evaluation results. Section VI compares with related work. Finally, Section VII outlines future work and concludes the paper.

II. Background

This section gives an overview on virtualization and memory management in cloud data centres, on Xen's transcendent memory and outlines the hardware architectures which implement coherence islands.

A. Virtualization Technology in the Cloud

There are many cloud service models available in the industry, such as SaaS (Software-as-a-Service), PaaS (Platform-as-a-Service) and IaaS (Infrastucture-as-a-Service). Considered one of the most important, IaaS allows users to dynamically access a configurable (seemingly unlimited) pool of computing resources. These resources are made available to the VMs using node-level virtualization software known as a hypervisor or Virtual Machine Manager (VMM), e.g. Xen [2], KVM [3] or VMWare [4].

1) Hypervisor and Dynamic Memory Management: The hypervisor virtualizes the computing resources of the node, and creates and manages Virtual Machines (VMs) with their own (guest) OSs. When a VM is created, the hypervisor allocates for it a portion of the physical memory. If the portion allocated is less than the amount of memory the VM needs at some point, then the VM is *under-provisioned* of memory. But when the VM has excess memory, the VM is *over-provisioned*. To optimize the utilization of memory and increase performance, memory has to be re-allocated from the VM that has an excess of memory to the VM that needs it.

State-of-the-art hypervisors implement mechanisms to dynamically re-allocate memory among VMs, including memory ballooning and memory hotplug. These mechanisms have been widely deployed in data centres, obtaining high levels of memory utilization. However, they do not provide adequate interfaces to aggregate memory capacity [22].

2) State-of-the-Art Transcendent Memory: Transcendent memory (Tmem) [1] was introduced as an additional solution for dynamic memory management. It pools the node's physical pages that are under-utilized and/or that are unassigned (fallow) to any VM. Tmem is abstracted as a key–value store in which pages are allocated/written and read through a paravirtualized put–get interface. Data written to Tmem is either *ephemeral* (e.g. clean data that the hypervisor can discard to reclaim memory) or *permanent* (e.g. dirty data that must be maintained) and either *private* (to a VM) or *shared* (among VMs).

Linux can take advantage of Tmem in two ways: *cleancache* and *frontswap*, which have been part of the kernel since Linux 3.0 (July 2011) and Linux 3.5 (July 2012), respectively. Both require a special Tmem kernel module to be loaded, in order to access the hypervisor's Tmem functionality using the relevant hypercalls. Linux cleancache is a cache for clean pages that are evicted by the Linux kernel's Pageframe Replacement Algorithm (PFRA). Linux frontswap uses Tmem as a swap device cache. When a page is getting swapped, frontswap attempts to store it (put) into Tmem, and it is swapped in case the put fails. Whenever the VM tries to access a page from the filesystem or swap space, it will check whether it was previously put into Tmem. If so, it will get it from Tmem assuming the page wasn't reclaimed. A successful put will avoid the disk write and read. Pages can also be de-allocated using flush-page operations, freeing the page to be used by other VM.


Fig. 1. UNIMEM architecture with two coherence islands

B. Hardware Support for Coherence Islands

Recent advances in computer architecture have proposed architectures with a shared global physical address space without global hardware cache coherence [6], [8], [10], [5], [7]. The essential features of architectures like EUROSERVER [7] are captured by the UNIMEM (*Unified Memory*) memory model.

UNIMEM is well suited for distributed systems that allow for resources to be shared. Figure 1 summarises how UNIMEM works. The number of nodes in current UNIMEM-based architectures ranges from 4 to 8 nodes. Each node contains N processor cores (ranging from 6 to 8 in current implementations), connected in clusters via a local cache-coherent interconnect to local DRAM and I/O devices. Remote memory is visible through the global physical address space, and communication among the nodes is achieved through an *inter-node interface* and *global interconnect* responsible for routing remote memory accesses to the appropriate node. We can summarize the essential characteristics of these architectures as:

- A global physical memory address space provides RDMA and load/store access to memory in all nodes.
- Routing uses the top bits of the global physical address.
- Fast low-latency communication is provided among the nodes of the system.
- Each node executes its own hypervisor and OSs.
- Features required for resiliency (see Section III-C4).

III. vMCA DESIGN

The vMCA stack consists of three components:

- Hypervisor support (Sec. III-A)
- Tmem Kernel Module (TKM) in Dom0 (Sec. III-B)
- Memory Manager (MM) in Dom0 (Sec. III-C)

A. Hypervisor Support

The Xen hypervisor has been extended to support vMCA, and all extensions are localized in the Tmem subsystem, which originally supports the necessary features to allocate and deallocate local Tmem pages on behalf of the VMs.

Page ownership: Each hypervisor *owns* a subset of the physical pages in the system, which constitutes the pool of memory that can be used for its own purposes and allocated to its VMs. When a hypervisor boots and first joins the vMCA system, it owns all of its *local* physical memory; i.e. all pages whose *home* is that node. If it needs additional memory capacity, the MMs may grant it ownership of additional *remote* memory pages. Conversely, it may release ownership of some

Memory Statistics	Description
$ENOMEM$	Code used in the hypervisor to signify that a *put* failed due to a lack of Tmem capacity
$node\_info$	Data structure that holds general status information of the computer host.
$node\_info.total\_tmem$	Total number of pages available to Tmem (free or allocated)
$node\_info.id$	Identifier of the node inside vMCA-supported system
vm_data_{hyp}	Data structure that holds the parameters of all of the VMs within the hypervisor
$vm\_data_{hyp}[id].vm\_id$	Identifier of the VM within Xen
$vm\_data_{hyp}[id].tmem\_used$	Number of Tmem pages currently used by the VM
$vm\_data_{hyp}[id].mm\_target$	Target number of pages for the VM, held by the hypervisor and previously sent by the MM
$vm\_data_{hyp}[id].puts\_total\_rate$	Total number of *puts* issued by the VM in the most recent period
$vm\_data_{hyp}[id].puts\_succ\_rate$	Total number of successful *puts* issued by the VM in the most recent period
$memstats$	Data structure that holds the last sampled statistics that the hypervisor sent to the MM.
$memstats.vm\_count$	Amount of active VMs as seen by the MM.
$memstats.vm$	Array where each entry holds statistics about an active VM
$memstats.vm[i].vm\_id$	Identifier of the VM within the MM, as received from the hypervisor
$memstats.vm[i].puts\_total\_rate$	Total number of *puts* that a VM has issued to the hypervisor in the recent period
$memstats.vm[i].puts\_succ\_rate$	Total number of successful *puts* that a VM has issued to the hypervisor in the recent period
$memstats.total\_alloc\_pages$	Total number of pages available in every zoned buddy allocator of the hypervisor
mm_out	Pointer to a data structure that holds the output parameters of the MM policy
$mm\_out[i].vm\_id$	VM identifier that maps a VM to its target allocation as calculated by the MM
$mm\_out[i].mm\_target$	Memory allocation target as calculated by the policy in the MM

TABLE I
SUMMARY OF MEMORY STATISTICS USED IN THE MEMORY MANAGER (MM)

Algorithm 1 Enforce memory allocation in the Hypervisor

```
 1: function DO_TMEM_PUT(vm_data_hyp, id)
 2:     tmem_used ← vm_data_hyp[id].tmem_used
 3:     mm_target ← vm_data_hyp[id].mm_target
 4:     if tmem_used ≥ mm_target then
 5:         return_value ← −ENOMEM
 6:     else if node_info.free_tmem == 0 then
 7:         return_value ← −ENOMEM
 8:     else
 9:         vm_data_hyp[id].tmem_used += 1
10:         vm_data_hyp[id].puts_succ += 1
11:         return_value ← 1
12:     end if
13:     vm_data_hyp[id].puts_total += 1
14:     return return_value
15: end function
```

of its local or remote pages, so that they can be granted to another node that has a greater need for them. As described in Section III-C, the MMs collectively ensure that each physical page in the system is owned by at most one hypervisor. The only times when a page is not owned by any hypervisor are when it is in transit between hypervisors, and, rarely, following a memory leak caused by a failed node (Section III-C4).

Page allocation: The pages owned by a hypervisor are allocated using a zoned Buddy allocator, with a zone for each node, including itself, from which it has ownership of at least one page. These allocators should be stored in a tree corresponding to the hierarchy of the system (nodes in leaves, then, e.g. boards, chasses and racks). A Tmem put traverses the tree to find the closest non-empty allocator, and allocates a page from it. A Tmem flush frees a page by returning it to the corresponding Buddy allocator.

Transfer of page ownership: In order to minimise communication overheads, page ownership is transferred between

hypercalls in *blocks*, each an appropriately-aligned power-of-two number of pages. The hypervisor supports four hypercalls in order to transfer page ownership. The Grant hypercall is used to receive ownership of a list of blocks. These blocks are added to the appropriate Buddy allocator. In contrast, a Request hypercall requests that the hypervisor release ownership of a given number of pages on behalf of another node. In response, the hypervisor takes as many blocks as needed from its allocators using a simple heuristic to prefer large blocks located physically close to the target node. The Return hypercall is issued whenever another node wishes to shut down or otherwise leave vMCA. It asks the hypervisor to release ownership of all pages whose home is that node; i.e. that are physically located on that node. It returns free pages, and, executing asynchronously, searches for pages in use by the Tmem clients and migrates their contents to free pages. Finally, the Invalidate-Xen hypercall is used if a remote node fails: it discards all free or allocated pages whose home is that node, and it terminates all VMs that were using such data.

Enforcing local per-VM memory constraints: The hypervisor enforces the Tmem allocations determined by the MM. Algorithm 1 illustrates this enforcing mechanism. The $do\_tmem\_put()$ function is called when a VM attempts to store a page in Tmem. The hypervisor then checks the amount of Tmem the VM has (line 4), and returns $-ENOMEM$, denying the request, if the VM is already at the limits of its allocation. If not (lines 9–11), the Tmem page will be given to the VM and the corresponding data structures will be updated.

Tmem Statistics: Table I presents the Tmem statistics gathered by the hypervisor. The statistics need to be minimum in order to minimise the communication overhead from the hypervisor to the MM. The hypervisor sends statistics to the MM every second.

Command	Direction	Description	Slave state
Distribution of global memory capacity			
Statistics(S)	S→M	Send node statistics S to Master	Active
Grant-Any(n, x)	S→M	Request n pages to slave x	Active
Grant-Fwd(n, x, y)	M→S	Forward request of n pages coming from y to slave x	Active
Force-Return(n, x)	M→S	Disable node and return all pages located at slave x (for leaving)	Active
Flow of page ownership			
Grant($b, \cdots$)	S/M→S	Transfer ownership of blocks of pages	Active
Node state changes			
Register	S→M	Register a new node	Inactive→Active
Leave-Req	S→M	Node requests to leave the system (e.g. for shutdown)	Active→Leaving
Leave-Notify	M→S	MM–M notifies that the recipient has left the system	Leaving→Inactive
Enable-Node(x, e)	M→S	Inform slave to accept ($e = 1$) or reject ($e = 0$) pages at node x	Active
Resiliency support			
Invalidate	S→M	Invalidate all pages located at sending node	Active/Leaving→Recovery
Invalidate(x)	M→S	Request recipient to invalidate pages at node x	Active
Invalidate-Notify(x)	M→S	MM–M notifies that the recipient has been invalidated	Recovery→Inactive

TABLE II

MEMORY MANAGER MESSAGE TYPES. IN THE TABLE, "S" STANDS FOR SLAVE AND "M" FOR MASTER

B. Dom0 Tmem Kernel Module (TKM)

The Tmem client interface requires a kernel module in each guest domain. vMCA needs a kernel module in the privileged domain Dom0 that acts as an interface between the hypervisor (using hypercalls) and the node's MM.

C. Dom0 User-space MM

Each node has a user-space MM in its privileged domain, Dom0. The MMs cooperate to:

- Distribute memory owned by each node among its guests
- Distribute global memory capacity among nodes
- Implement the flow of page ownership among nodes
- Enable nodes to join and leave vMCA, and handle failures

In the current design, one of the MMs is designated to be the Memory Manager Master (MM–M), which is responsible for system control and global memory capacity distribution. The MMs communicate using a secure and reliable packet transport such as SSL/TLS. The messages passed among the MMs are listed in Table II.

1) Joining the vMCA system: In order to join vMCA, a node requires a configuration file, which provides the network address of the nodes, security credentials to establish secure connections, and the mapping for all nodes from node ID to network address. For managing locality, it also needs to know the location of each node in the NUMA hierarchy.

When a node R wishes to join vMCA, it first sends a Register message to the MM–M (see Table II). The MM–M then sets its state to Active (Figure 2) and sends an Enable-Node(R,1) message to all of the registered nodes. Each node maintains a bitmap of the enabled nodes.

2) Distributing ownership of memory: The flow of memory pages across vMCA considers many aspects.

Distributing memory owned by a node among guests: Each node's MM determines the maximum Tmem allocation for each VM, using a policy (Section III-D) that exploits the Tmem statistics gathered by the hypervisor. Pages are dynamically allocated to every VM subject to total (local plus remote) consumption limits, which are sent to the hypervisor using a Mem-Limit hypercall.

For every policy, the following condition must be met:

$$\sum_{i=1}^{i=m} vm\_data_{\text{MM}}[i].mm\_target = total\_tmem(t_i) \quad (1)$$

Eq. 1 means that at time t_i, the amount of total Tmem pages available to the node ($total\_tmem$) is equal to the sum of the allocations of all m active VMs. This implies that every Tmem page is usable, and that Tmem pages are not overallocated. However, when calculating the allocation limits for every VM, it is possible that the sum of the allocations exceeds $total\_tmem(t_i)$.

Overallocation of pages generates a situation in which one or more VMs take memory in excess, potentially preventing all VMs from reaching their target allocations, defeating the purpose of the management policies. When over-allocation occurs, the MM recalculates the target of the VMs according to:

$$tgt_{vm_i} = \frac{total\_tmem \times vm\_data_{\text{MM}}[i].mm\_target}{\sum_{i=1}^{i=m} vm\_data_{\text{MM}}[i].mm\_target} \quad (2)$$

Eq. 2, where tgt_{vm_i} is an abbreviated form of $vm\_data_{\text{MM}}[i].mm\_target$, ensures that the proportional allocation of each VM follows the policy's desired proportion while also satisfying Eq. 1.

Distributing global memory capacity among nodes: All nodes in the Active state send statistics to the MM–M using the Statistics message. A node that needs memory issues a Grant-Any message to the MM–M , and the latter redistributes the memory based on the statistics and the global memory policy. The MM–M redistributes the memory capacity among nodes by sending Grant-Fwd messages, which forwards requests to a donor node to transfer ownership of a specified number of free physical pages to a requesting node.

Flow of page ownership: The MM–M rebalances memory capacity without knowing the actual physical addresses.

Fig. 2. Node state transition diagram (at MM–M)

Ownership of global physical addresses is transferred peer-to-peer using Grant messages. A Grant message passes a list of blocks, each an appropriately-aligned power-of-two number of physical pages. The MM–M itself can sometimes satisfy the request for memory of another node, for which it will return a Grant message to the requesting node.

To avoid race conditions during node shutdown and failure, the recipient checks the home of each received block against the bitmap of enabled nodes (Section III-C1). If the home node of a block is disabled (rarely occurs), then ownership is returned back to its home using a new point-to-point Grant.

3) Leaving the vMCA system: To shutdown a node of vMCA, the following procedure must be followed. Note that if a node fails, then the procedure in Section III-C4 is followed.

1) The node R that wishes to shutdown sends Leave-Req to the MM–M.
2) In response, the MM–M moves the node to the Leaving state (see Figure 2). It also sends a Force-Return(R) message to all nodes in the system. Each recipient returns all of the pages at R that it owns and will reject any such pages received in future Grant messages.
3) Node R frees all pages used by Tmem and returns ownership of all remote pages to their home nodes using peer-to-peer Grant messages.
4) Periodically, each node sends a Grant message to node R to return ownership of the pages that it had borrowed.
5) Once the MM–M has received Statistics from every node indicating that R is disabled and that it owns no pages at R, then the MM–M moves R to Inactive (Figure 2) and sends Leave-Notify to R.
6) At this point the node R has ownership of all its non-leaked pages, has left vMCA, and may shutdown.

4) Resilient memory capacity aggregation: There are two aspects to ensuring resiliency in the face of node failures: 1) protecting the integrity of the data stored in Tmem, and 2) restoring lost system state.

Ensuring data integrity on slave failure: The Tmem interface guarantees a get returns exactly the data previously written by the corresponding put. In addition, if the pool is *permanent*, the get operation must succeed.

Assume that node A has ownership of memory located at node H. If H fails, then this data is lost. If H reboots, it will restart and may even attempt to re-join vMCA. Since any get performed at A must never access the (incorrect) new contents of the physical page at node H, the following procedure must be followed after a failure:

1) When a node R boots, remote memory accesses to it must raise a hardware exception (Section III-E), which is also raised when a node attempts to access a page from a failed node. If any remote page access causes such an exception, then if the page is in *permanent* pool, the corresponding VM must be shutdown. If the page is *ephemeral*, the get fails.
2) When the MM at node R is initialized without a previous clean shutdown, it first sends Invalidate to the MM.
3) The MM–M moves the node to the Recovery state (Figure 2) and sends Invalidate(R) to all nodes. Each node then disables node R, and if it has ownership of pages with home R, it makes an Invalidate-Xen hypercall to invalidate the pages.
4) Once the MM–M has received Statistics from every node, i.e. it has disabled R and owns no pages at R, then the MM–M moves the node to the Inactive state and sends an Invalidate-Notify message to R.
5) Upon receiving the Invalidate-Notify message, node R begins its normal initialization procedure.

Restoring lost pages on slave failure: Following a failure of a given node R, vMCA is unable to recover the pages that were owned by R. It is possible to approximate the number of lost pages from a home node using the statistics sent by the nodes, by adding the number of pages owned and comparing against the number of physical pages at the node. However, since this is done using potentially outdated information from Statistics messages, the value will not be exact.

Restoring system state on MM–M failure: If the MM–M fails, the system will continue to work except that a) no memory capacity will be redistributed, and b) no node can enter/leave the system.

When the MM–M attempts to restart, it will listen for connections from the other nodes. When a node connects and starts sending Statistics, it is added to the system in the Active state. A node that is in the process of leaving or invalidation will send a new Leave-Req or Invalidation message on re-connecting to the MM–M, which will restart the leaving or invalidation process. At any time, all nodes in the Active state are part of the global memory allocation.

D. Memory Management Policies

We implemented two memory management policies: 1) greedy-remote [9] and a 2) two-level memory management policy (TLP), divided in a first and a second level of memory management. We also test a third one: greedy-local [9] which is the baseline single-node Tmem implementation.

The case of greedy-remote allows for the node's hypervisor to give Tmem pages away to its VMs on demand without any constraints (similar to greedy-local). It differs from greedy-local in that whenever it detects that a node is under pressure, greedy-remote issues a Grant-Any to the MM–M requesting for fixed number (1000) of pages.

Algorithm 2 First-Level of TLP

1: ▷ *FLM is called in all nodes when the MM receives statistics from the hypervisor (every second).*
2: **function** FLM_POLICY($memstats, node\_info, P, threshold$)
3:     $ttmem \leftarrow memstats.total\_alloc\_pages$
4:     $node\_id \leftarrow node\_info.id$
5:     $sumtgt \leftarrow 0$
6:     **for** $i \leftarrow 1, memstats.vm\_count$ **do**
7:         $puts\_total\_rate \leftarrow memstats.vm[i].puts\_total\_rate$
8:         $puts\_succ\_rate \leftarrow memstats.vm[i].puts\_succ\_rate$
9:         $puts\_fail\_rate \leftarrow puts\_total\_rate - put\_succ\_rate$
10:         **if** $puts\_fail\_rate > 0$ **then**
11:             $ctgt \leftarrow memstats.vm[i].mm\_target$
12:             $mm\_target \leftarrow ctgt + (P * ttmem)/100$
13:         **else**
14:             $ctgt \leftarrow memstats.vm[i].mm\_target$
15:             $curr\_use \leftarrow memstats.vm[i].tmem\_used$
16:             $difference \leftarrow ctgt - curr\_use$
17:             **if** $difference > threshold$ **then**
18:                 $mm\_target \leftarrow ((100 - P) * ctgt)/100$
19:             **else**
20:                 $mm\_target \leftarrow ctgt$
21:             **end if**
22:         **end if**
23:         $mm\_out[i].vm\_id \leftarrow memstats.vm[i].vm\_id$
24:         $mm\_out[i].mm\_target \leftarrow mm\_target$
25:         $sumtgt \leftarrow sumtgt + mm\_target$
26:     **end for**
27:     **if** $sumtgt > ttmem$ **then**
28:         $send\_msg($Grant-Any$(node\_id, (sumtgt - ttmem)))$
29:         **for** $i \leftarrow 1, memstats.vm\_count$ **do**
30:             $new \leftarrow (ttmem/sumtgt) * mm\_out[i].mm\_target$
31:             $mm\_out[i].mm\_target \leftarrow new$
32:         **end for**
33:     **end if**
34:     $send\_hypercall($Mem-Limit$(mm\_out))$
35:     $send\_msg($Statistics$(memstats))$
36: **end function**

Algorithm 3 Second-level of TLP

1: ▷ *SLM is called when a* Grant-Any *(GrAny) message is received.*
2: **function** SLM($memstats, GrAny, R, np \leftarrow \{\}$)
3:     **for** $i \leftarrow 1, memstats.node\_count$ **do**
4:         $np[i].id \leftarrow memstats[i].node\_id$
5:         $np[i].tmem \leftarrow memstats[i].total\_tmem$
6:     **end for**
7:     $sort\_nodes(np)$
8:     **for** $i \leftarrow 1, memstats.node\_count \wedge GrAny.n > 0$ **do**
9:         $diff \leftarrow np[i].tmem - GrAny.n$
10:         **if** $diff > R$ **then**
11:             $nr \leftarrow GrAny.n$
12:             $send\_msg($Grant-Fwd$(nr, np[i].id, GrAny.x))$
13:             $GrAny.n \leftarrow 0$
14:         **else if** $diff < R \wedge np[i].tmem \geq R$ **then**
15:             $nr \leftarrow np[i].tmem - R$
16:             $send\_msg($Grant-Fwd$(nr, np[i].id, GrAny.x))$
17:             $GrAny.n \leftarrow GrAny.n - nr$
18:         **end if**
19:     **end for**
20: **end function**

Algorithm 2 shows the first level of *TLP*. It uses the statistics provided by the hypervisor ($memstats$, in Table I), checks the number of VMs running (line 6) and if the VMs have failed $puts$ since receiving the previous statistics set (lines 7–9). Initially, all VMs have an equal portion of the node's available Tmem. If $put\_fail\_rate > 0$, the MM increases the allocation of the VM by a percentage P of the total amount of Tmem pages available to the node (lines 10–13).

If $put\_fail\_rate = 0$ (it can't be negative), the VM allocation is reduced by a percentage P assuming that it currently has more pages allocated than its current use by a $threshold$ value ($difference > threshold$). In case not, then it keeps the current allocation of the VM (lines 13–22).

The MM then sums and assigns the target allocations to the corresponding data structures (lines 23–25). If the sum exceeds the amount of Tmem available to the node, the MM does the following: 1) send Grant-Any to the MM–M to request memory pages (line 28), and 2) readjust the allocation of the VMs according to equation 2 (lines 29–32). The MM issues the Grant-Any because overallocation implies that the node is under memory pressure, so it needs remote pages. There's no guarantee it will get them after requesting it, thus it's necessary

to readjust the allocations until more pages become available. Once all necessary calculations have been done, the MM issues the Mem-Limit hypercall (line 34) and the Statistics message to the MM–M if the node is a slave (line 35).

Upon receiving the Grant-Any request, the MM–M decides how many pages to actually request and from which node(s) to get them from. This is the second-level of TLP shown in Algorithm 3. The MM–M has an extended version of the $memstats$, where it stores the statistics from itself and from other slave nodes that it gets through Statistics messages.

Algorithm 3 uses a data structure (np) (line 2) that stores the amount of Tmem pages available in each node and their IDs (lines 3–6). Then, it sorts the nodes (using merge sort) in descending order in np (line 7). This is to decide which node to forward the request to (lines 8–18), usually the one with more memory available. The MM–M first checks that the potential donor node will have enough memory for itself (above a threshold R). If the node is able to satisfy the request while meeting the threshold, then the request is forwarded to it using a Grant-Fwd message (lines 9–14).

If the node is unable to satisfy the request but has memory available, the MM–M can still forward the request (line 16) but requesting less pages (line 15) to meet the threshold of the node. Then, a new node from np is selected to request the rest of the remaining pages. From this process, it's clear that there's no guarantee that requests for pages will be satisfied.

E. Hardware Support for Memory Aggregation

vMCA requires hardware with the following features:

1) Fast interconnect providing a synchronous interface to the NUMA distributed memory.
2) Direct memory access from the hypervisor to all the memory available: the hypervisor has to be able to access transparently its local and remote memory through the same mechanisms, either through load/store instructions or through RDMA.

Node	CPU	Frequency	Memory
Node 1	AMD FX Quad-Core	1.4 GHz	8 GB
Node 2	Intel Core i7	2.10 GHz	16 GB
Node 3	Intel Xeon	2.262 GHz	64 GB
Node 4	AMD FX-4300 Quad-Core	3.8 GHz	8 GB

TABLE III
HARDWARE CHARACTERISTICS

3) Remote access to a node's pages is disabled on hardware boot. Access is enabled when the node joins vMCA, when it sends the Register message.
4) Extraction of the node ID given a physical address, for instance depending on some higher-order bits.

IV. EXPERIMENTAL METHODOLOGY

We evaluated vMCA in a platform consisting of four computing nodes. UNIMEM-based architectures usually have 4 to 8 nodes sharing the global physical address space. Every node and VM runs Ubuntu 14.04 with Linux kernel 3.19.0+ as the OS, and Xen 4.5. The MM of the nodes communicate using TCP/IP sockets over Ethernet. Node 2 executes the MM–M. The hardware properties of the nodes are given in Table III.

The shared global address space is emulated using the node's local memory. Xen has been modified to start using a portion of the physical memory capacity, equalling the emulated memory capacity of the node. The rest of the node's memory capacity was reserved to emulate remote data storage. Whenever the hypervisor performs an "emulated" remote access, we add a delay in the hypervisor lasting 50 μs to model hardware latency. We evaluate vMCA using the CloudSuite 3.0 Benchmarks [11], with every VM having 1 vCPU and every node having 1 GB of Tmem initially available. We execute at most three VMs simultaneously, and refer to each set of VMs as a *scenario*. Table IV shows our scenarios.

V. RESULTS

A. Results for Scenario 1

Figures 3(a) and 3(b) present the average running times of each VM for Scenario 1 for the three policies. They show average improvements of 11.94% and 12.9% in nodes 3 and 4, respectively, when going from greedy-local to greedy-remote. With TLP, there is a further improvement of 7.3% (TLP, $P = 2.0\%$) and 7.0% (TLP, $P = 6.0\%$) in nodes 3 and 4 over greedy-remote.

Figures 3(c), 3(d), and 3(e) present the amount of Tmem (remote and local) that each VM takes for the three policies in node 3. For greedy-local and greedy-remote shown in Figures 3(c) and 3(d), VM3 and VM4 take a smaller proportion of Tmem when compared to VM2, but this disparity reduces with TLP, as shown in Figure 3(e). TLP tries to be *fair* on how the Tmem pages are distributed among the VMs.

For some values of P, TLP degrades the performance of in-memory-analytics. Consider the first iteration of VM1 in node 3 for $P = 1$ (tlp-1, Figure 3(a)). With TLP, limits are enforced on the amount of Tmem a VM can have. When in-memory-analytics tries to go beyond its initial allocation,
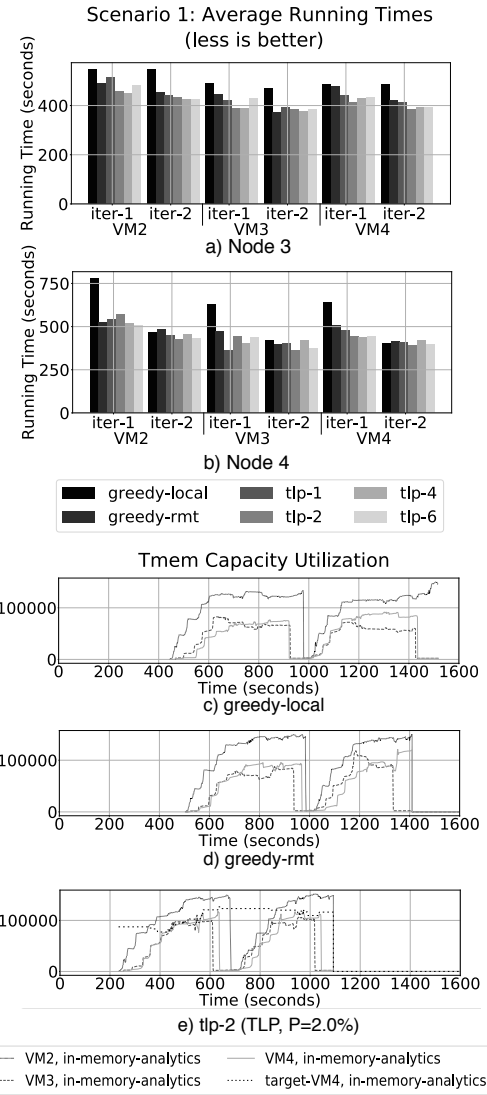

Fig. 3. Running time for Scenario 1 in nodes 3 (a) and 4 (b). Utilization of the Tmem capacity (nod-tmem) by every VM in node 3 for Scenario 1 in all three policies (c, d ,e). *target-VM4* refers to the target allocation of VM4.

the hypervisor prevents the VM from taking pages until the MM updates its allocation target. But with $P = 1$, the targets increase at a slower pace than needed. Thus, the VM generates disk accesses, degrading its performance.

With higher P, the MM increases the targets faster. But when P becomes too high, a VM might have excess memory allocated, making it unlikely for other VMs to take a fair share, thus degrading performance. In Figures 3(a), the performance improves when increasing P up to 2.0% but degrades again as P keeps increasing. This highlights a trade-off between the adaptability of TLP to adjust to changes in memory demand and to achieve a fair distribution of Tmem.

B. Results for Scenario 2

Figures 4(a) and 4(b) present the average running times of each VM for Scenario 2 for the three policies. They show average improvements of 50.84% and 45.96% in nodes 3 and 4, respectively, when going from greedy-local to greedy-

| Scenario	VM RAM (MB)	Description
Scenario 1	VM1: 768, VM2: 1024, VM3: 1024	Every VM executes simultaneously in-memory-analytics once, sleeps 5 seconds and then executes it again. The data set was taken from [20]
Scenario 2	VM1: 768, VM2: 768, VM3: 768	Every VM executes simultaneously graph-analytics once. The data set was taken from [19], [17], [18]
Scenario 3	VM1: 768, VM2: 768, VM3: 1024	VM1 and VM2 execute graph-analytics while VM3 executes in-memory-analytics, all at once. Every VM executes its benchmarks twice.
Scenario 4	VM1: 768, VM2: 512, VM3: 512	VM1 executes graphs-analytics, while VM2 and VM3 execute the client and server (memcached [16]) for data-caching, respectively.

TABLE IV
LIST OF SCENARIOS USED FOR BENCHMARKING
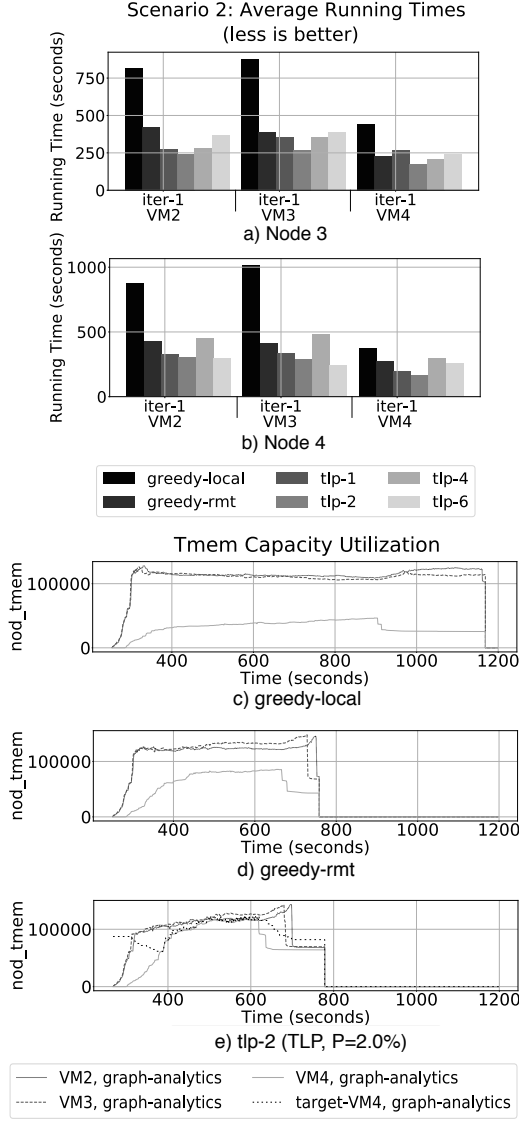
Fig. 4. Running time for Scenario 2 in nodes 3 (a) and 4 (b). Utilization of the Tmem capacity (nod-tmem) by every VM in node 3 for Scenario 2 in all three policies (c, d ,e). *target-VM4* refers to the target allocation of VM4.

remote. With TLP, there is a further improvement of 32.1% (TLP, $P = 2.0\%$) and 31.8% (TLP, $P = 2.0\%$) in nodes 3 and 4, respectively, over greedy-remote.

Figures 4(c), 4(d) and 4(e) show the amount of Tmem taken by each VM on node 3, for each policy. When $P = 2.0\%$, the

MM enforces *fairness* in the Tmem allocation of the VMs when compared to greedy-local and greedy-remote. The difference in Tmem ownership between VM4 and other VMs is massive for greedy-local but reduces for greedy-remote, mainly because of the availability of the remote memory. In both cases, the VMs are competing for the available Tmem.

Figures 4(a) and 4(b) highlight again the *adaptability-vs-fairness* tradeoff. As P increases, the performance improvements hits a minimum and then increases. For this scenario, this trend is maintained for larger values of P, making $P = 2.0\%$ an optimal value in this case.

C. Results for Scenario 3

Figures 5(a) and 5(b) present the average running times of each VM for Scenario 3 for the three policies. They show average improvements of 50.6% and 55.4% in nodes 3 and 4, respectively, when going from greedy-local to greedy-remote. When implementing TLP, there is a further improvement of 10.1% (TLP, $P = 6.0\%$) and 15.6% (TLP, $P = 4.0\%$) in nodes 3 and 4 over greedy-remote.

Figures 5(c), 5(d) and 5(e) show the amount of Tmem that each VM takes for the three policies in node 3. As remote memory becomes available to the node, the running time dramatically drops. VM4, running in-memory-analytics, takes the longest due to the nature of the benchmark. With TLP, the amount of Tmem that VM2 and VM3 have become similar for the second iteration of graph-analytics achieving fairness, while they present an important difference for the second iteration for greedy-local and greedy-remote.

The adaptability-vs-fairness tradeoff is observed in the running times of Figures 5(a) and 5(b). Performance improves as P increases, but in some cases it changes suddenly when P changes, like in the second iteration of VM1 and VM2 for node 3 (Figure 5(a)). In both nodes, the second iteration of VM4 (in-memory-analytics) performs the same for all policies because it runs on its own, while the other two VMs have completed (Figure 5(c,d,e)). In this Scenario, it is more difficult to obtain an optimal value of P.

After VM2 and VM3 execute graph-analytics, they keep a large part of the Tmem pages they used. The MM is unable to reclaim Tmem pages that are still *in use*, only becoming available when the VMs explicitly issue flush operations. But the VMs never flush, thus keeping the pages. If the node runs more VMs with benchmarks that interact with Tmem in this way, the Tmem pages may be depleted inefficiently over time.
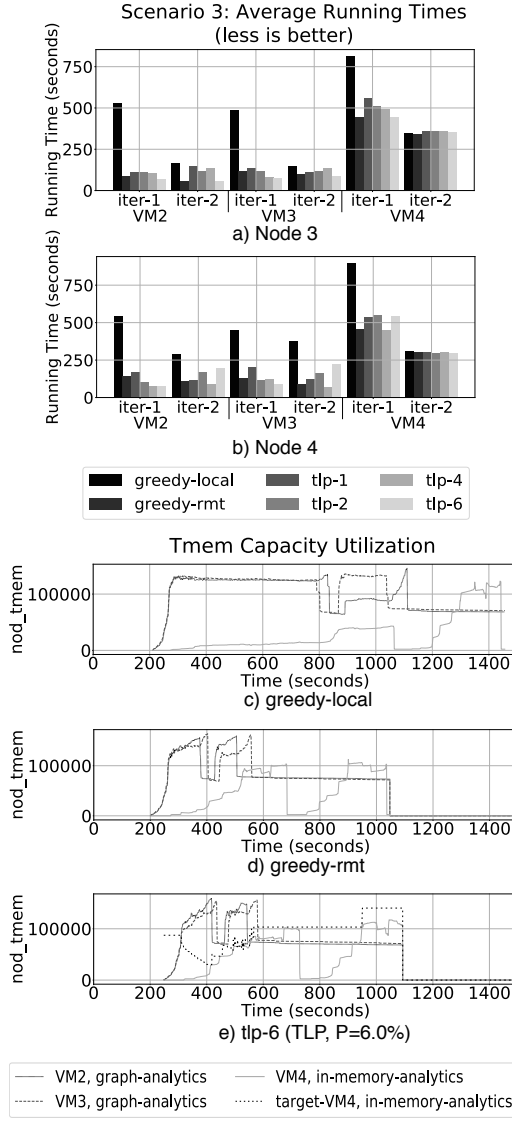
Fig. 5. Running time for Scenario 3 in nodes 3 (a) and 4 (b). Utilization of the Tmem capacity (nod-tmem) by every VM in node 3 for Scenario 3 in all three policies (c, d ,e). *target-VM4* refers to the target allocation of VM4.

Fig. 6. Running time for Scenario 4 in nodes 3 (a) and 4 (b). Utilization of the Tmem capacity (nod-tmem) by every VM in node 3 for Scenario 4 in all three policies (c, d ,e). *target-VM4* refers to the target allocation of VM4.

D. Results for Scenario 4

Figures 6(a) and 6(b) presents the average running times of each VM for Scenario 4 for the three policies. They show average improvements of 63.5% and 76.8% in node 3 and node 4, respectively, when going from greedy-local to greedy-remote. With TLP, it shows improvements of 23.7% (TLP, $P = 2.0\%$) and 37.5% (TLP, $P = 2.0\%$) in nodes 3 and 4 over greedy-remote.

Figures 6(c), 6(d) and 6(e) show the total Tmem capacity of each VM for the three policies in node 3. In all cases, VM4 (memcached server) uses the same amount of Tmem, while VM2 (graph-analytics) and VM3 (data-caching-client) compete for the available Tmem. With TLP (Figure 6(e)), fairness is achieved fairly quickly and the adaptability-vs-fairness tradeoff is evident in Figures 6(a) and 6(b). In this case, $P = 2.0\%$ seems to be optimal for most VMs, except for VM3 in node 4, in which $P = 1.0\%$ performs better.
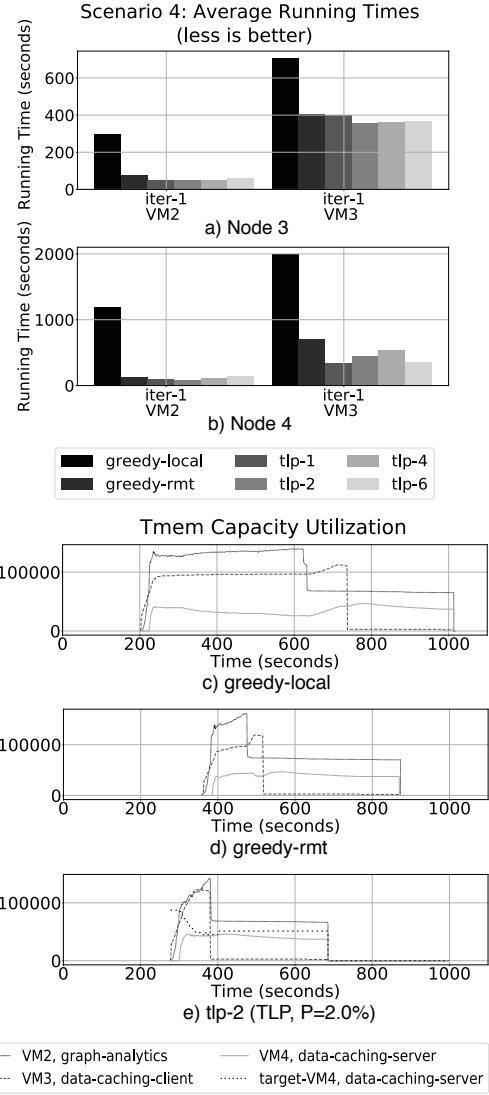
VI. RELATED WORK

Venkatesan et al. [21] use Tmem and non-volatile memory (NVM) to reduce accesses to disk by VMs in a system with DRAM and NVM. This work is relevant to use because it shows that Tmem is useful in systems with non-uniform memory hierarchies. The non-uniformity in our case is a consequence of the different memory technologies across the servers, the amount of memory available to any given server and the physical location of this memory.

Lorrillere et al. [22] proposed a mechanism, called PUMA, to improve memory utilization based on remote caching to pool VMs memory across a data center. PUMA pools VMs' unused memory for the benefit of other VMs having I/O intensive workloads. PUMA makes use of an interface very similar to Tmem to access the cache present in a remote server. vMCA differs from [22] because the latter targets clean file-backed pages for I/O intensive applications, while vMCA

targets pages generated by the processes of the applications during their computations, which are not file-backed.

Zcache [13] is a backend for frontswap and cleancache that provides a compressed cache for swap and clean filesystem pages. *RAMster* [13] is an extension of zcache that uses kernel sockets to store pages in the RAM of remote nodes. Similarly, RAMCloud [14] is a POSIX-like filesystem in which all data is stored in RAM across the nodes. vMCA differs from all these approaches in: 1) vMCA grants and releases memory capacity at a higher granularity larger than a single page (reducing communication between nodes), 2) vMCA is not entirely within the kernel, leveraging user-space flexibility where possible, and 3) vMCA exploits a global shared address space with low-latency communication mechanisms.

Hwang et al. [23] proposed Mortar: a mechanism that pools spare memory on a server and exposes it as a volatile data cache. This data cache is managed by the hypervisor and it can evict objects from the cache in order to reclaim memory for other VMs. They test two uses cases, the first of which uses memcached protocol to aggregate free memory across data center, and the second works at the OS-level to cache and prefetch disk blocks. They also use a very similar interface to that provided by Tmem. vMCA puts all the memory aggregation and management details on a user-space process, while keeping the hypervisor small and secure.

VII. Conclusions and Future Work

This paper introduces vMCA, a resilient mechanism that exploits Tmem to aggregate memory capacity across multiple nodes. We evaluated vMCA using CloudSuite, obtaining up to 37.5% performance improvement when enabling memory aggregation and management policies. The results demonstrate the effectiveness of vMCA for improving the performance of the evaluated Cloudsuite applications.

Future work will investigate how to integrate our mechanism with other resource management mechanisms and extend vMCA to create a more complete solution for data centers. For example, we need to consider how feasible it is to reclaim Tmem pages in use and store them in disk, consider VM migration and ways for nodes to recover the memory they borrowed. It is desirable also to implement memory management policies that anticipate changes in memory demand.

References

[1] Magenheimer, D., Mason, C., McCracken, D., Hackel, K.: Transcendent memory and Linux. In: Proceedings of the Linux Symposium, pp. 191–200. Citeseer, 2009.

[2] Barham, P., Dragovic, B., Fraser, K., Hand, S., Harris, T., Ho, A., Neugebauer, R., Pratt, I., Warfield, A.: Xen and the art of virtualization. In: ACM SIGOPS Operating Systems Review, vol. 37, pp. 164–177. ACM, 2003.

[3] Kivity, A., Kamay, Y., Laor, D., Lublin, U., Liguori, A.: KVM: the Linux virtual machine monitor. In: Proceedings of the Linux symposium, pp. 225–230. 2007

[4] VMware Virtualization for Desktop & Server, Application, Public & Hybrid Clouds http://www.vmware.com

[5] Dong, J., Hou, R., Huang, M., Jiang, T., Zhao, B., Mckee, S., Wang, H., Cui, X., Zhang, L.: Venice: Exploring Server Architectures for Effective Resource Sharing. In: IEEE International Symposium on High-Performance Computer Architecture (HPCA), 2016.

[6] Howard, J., Dighe, S., Hoskote, Y., Vangal, S., Finan, D., Ruhl, G., Jenkins, D., Wilson, H., Borkar, N., Schrom, G.: A 48-core IA-32 message-passing processor with DVFS in 45nm CMOS. In: International Solid-State Circuits Conference Digest of Technical Papers (ISSCC), pp. 108–109. IEEE (2010).

[7] Durand, Y., Carpenter, P., Adami, S., Bilas, A., Dutoit, D., Farcy, A., Gaydadjiev, G., Goodacre, J., Katevenis, M., Marazakis, M., Matus, E., Mavroidis, I., Thomson, J.: Euroserver: Energy Efficient Node For European Micro-servers. In: 17th Euromicro Conference on Digital System Design (DSD), pp. 206–2013. IEEE (2014).

[8] Katrinis, K., Syrivelis, D., Pnevmatikatos, D., Zervas, G., Theodoropoulos, D., Koutsopoulos, I., Hasharoni, K., Raho, D., Pinto, C., Espina, F., Lopez-Buedo, S., Chen, Q., Nemirovsky, M., Roca, D., Klosx, H., Berends, T.: Rack-scale Disaggregated cloud data centers: The dReDBox project vision. In: Design, Automation & Test in Europe Conference & Exhibition (DATE). IEEE (2016).

[9] Garrido, L. A., Carpenter, P.: Aggregating and Managing Memory Across Computing Nodes in Cloud Environments. In: Proceedings of 12th Workshop on Virtualization in High-Performance Cloud Computing. Accepted paper.

[10] Thacker, C.: Beehive: A many-core computer for FPGAs. In: MSR Silicon Valley. 2010.

[11] Ferdman, M., Adileh, A., Kocberber, O., Volos, S., Alisafaee, M., Jevdjic, D., Kaynak, C., Popescu, A. D., Ailamaki, A., Falsafi, B.: Clearing the clouds: a study of emerging scale-out workloads on modern hardware. In: Proceedings of the seventeenth international conference on Architectural Support for Programming Languages and Operating Systems, pp. 37–48. ACM, London, England (2012).

[12] Rigo, A. et al. Paving the way towards a highly energy-efficient and highly integrated compute node for the Exascale revolution: the ExaNoDe approach. Accepted to Euromicro DSD 2017.

[13] Magenheimer. D.: Zcache and RAMster (oh, and frontswap too) overview and some benchmarking. https://oss.oracle.com/projects/tmem/dist/documentation/presentations/LSFMM12-zcache-final.pdf. 2012

[14] Ousterhout, J., Agrawal, P., Erickson, D., Kozyrakis, C., Leverich, K., Mazières, D., Mitra, S., Narayanan, A., Parulkar, G., Rosenblum, M., Rumble, S., Stratmann, E., Stutsman, R.: The case for RAMClouds: scalable high-performance storage entirely in DRAM. In: SIGOPS Operating Systems Review, vol. 43, pp. 92–105. ACM (2010).

[15] Svärd, P., Hudzia, B., Tordsson, J., Elmroth, E.: Hecatonchire: Towards Multi-host Virtual Machines by Server Disaggregation. In: Euro-Par International Parallel Processing Workshops, pp. 519–529. Springer International Publishing, Portugal (2014).

[16] memcached - a distributed memory object caching system. http://memcached.org

[17] Ross, R.A., Ahmed, N.K.: The Network Data Repository with Interactive Graph Analytics and Visualization. In: Proc. of the 29th AAAI Conference on AI. 2015.

[18] Ross, R.A., Ahmed, N.K.: An Interactive Data Repository with Visual Analytics. In: SIGKDD Explor., vol 17, no. 2, pp. 37–41. 2016.

[19] Cho, E., Myers, S.A., Leskovec, J.: Friendship and mobility: user movement in location-based social networks. In: Proc. of the 17th ACM SIGKDD International conference on Knowledge discovery and data mining, pp. 1082–1090. ACM. 2011.

[20] Harper, F. M., Konstan, J.A.: The MovieLens Datasets: History and Context. In: ACM Trans. on Interactive Intelligent Systems, pp. 19:1-19:19. ACM, USA (2015).

[21] Venkatesan, V., Qinqsong, W., Tay, Y.C.: Ex-Tmem: Extending Transcendent Memory with Non-volatile Memory for Virtual Machines. In Proceedings of the 2014 IEEE Intl Conf on High Performance Computing and Communications, 2014 IEEE 6th Intl Symp on Cyberspace Safety and Security, 2014 IEEE 11th Intl Conf on Embedded Software and Syst (HPCC,CSS,ICESS), pp. 966–973. Washington, DC, USA. 2014.

[22] Lorrillere, M., Sopena, J., Monnet, S., Sens, P.: Puma: Pooling Unused Memory in Virtual Machines for I/O Intensive Applications. In Proceedings of the 8th ACM International Systems and Storage Conference, pp. 1:1–1:11. Haifa, Israel, 2015.

[23] Hwang, J., Uppal, A., Wood, T., Huang, H.: Mortar: Filling the Gaps in Data Center Memory. In Proceedings of the 10th ACM SIGPLAN/SIGOPS International Conference on Virtual Execution Environments, pp. 53–64. New York, USA, 2014.